

# Assignment 3 Deep Generative Models

Qiao Ren

UvAnetID: 11828668

university of amsterdam

Dec, 2020

## Question 1.1

Step1 draw samples  $z_n$  from normal distribution:  $p(z_n) = N(0, I_D)$

Step2 compute:  $decoder(z_n) = f_\theta(z_n)$ . the result of decoder will be the mean of the probability distribution of each pixel, in the next step.

Step3 Each pixel is drawn from a bernoulli distribution. The mean of the bernoulli distribution is  $f_\theta(z_n)_m$ . So the probability distribution of a single pixel is:

$$p(x_n^{(m)}) = Bern(x_n^{(m)} | f_\theta(z_n)_m)$$

Step4 For the  $n^{th}$  img, it has  $M$  pixels in total. So the probability distribution of the  $n^{th}$  img is the joint probability of all the pixels.

$$p(x_n | z_n) = \prod_{m=1}^M Bern(x_n^{(m)} | f_\theta(z_n)_m)$$

## Question 1.2

Monte-Carlo integration is inefficient. Because:

1) As shown in Figure 2, the region of  $p(z|x)$  only covers a small part of the region  $p(z)$ . When we draw samples  $z_i$  from  $p(z)$ , there are lots of samples in the region of  $p(z)$  but out of the region of  $p(z|x)$ . For these samples,  $p(x_n|z_i)$  is almost zero. Therefore they don't contribute to computing  $\log p(x_n)$ . It means that we draw a large amount of useless samples  $z_i$ .

2) Monte-Carlo integration also suffers from the curse of dimensionality. When the dimensionality of  $z$  is high, each  $z_i$  is high dimensional. The computation cost is huge. We have to draw a huge amount of samples to cover the whole space of  $z$ . However it is infeasible to do so.

## Question 1.3

Example 1	$q = N(\mu_q, \sigma_q^2) = N(1, 10)$ $p = N(\mu_p, \sigma_p^2) = N(1.001, 10.001)$	$D_{KL}(q \parallel p)$ is very small
Example 2	$q = N(\mu_q, \sigma_q^2) = N(1, 10)$ $p = N(\mu_p, \sigma_p^2) = N(8000, 8000)$	$D_{KL}(q \parallel p)$ is very large

#### Question 1.4

$$\log p(x_n) - KL(q(Z | x_n) \parallel p(Z | x_n)) = ELBO$$

$$\log p(x_n) - KL(q(Z | x_n) \parallel p(Z | x_n)) = E_{q(z_n | x_n)}[\log p(x_n | z_n)] - KL(q(z | x_n) \parallel p(z))$$

We have computed that:  $\log(p(x_n)) = \log(E[\frac{p(z_n)}{q(z_n | x_n)} p(x_n | z_n)])$

We apply the Jensen's Inequality. Because log is a concave function. So it satisfies Jensen's Inequality.

$$h(E[function]) \geq E[h(function)]$$

$$\Rightarrow \log(E[\frac{p(z_n)}{q(z_n | x_n)} p(x_n | z_n)]) \geq E[\log(\frac{p(z_n)}{q(z_n | x_n)} p(x_n | z_n))]$$

$$\Rightarrow \log(E[\frac{p(z_n)}{q(z_n | x_n)} p(x_n | z_n)]) \geq E[\log(p(x_n | z_n))] + E[\log(\frac{p(z_n)}{q(z_n | x_n)})]$$

$$\Rightarrow \log(E[\frac{p(z_n)}{q(z_n | x_n)} p(x_n | z_n)]) \geq E[\log(p(x_n | z_n)) - KL(q(z | x_n) \parallel p(z))]$$

$$\Rightarrow \log(p(x_n)) \geq E[\log(p(x_n | z_n)) - KL(q(z | x_n) \parallel p(z))]$$

We must optimize the lower bound, instead of optimizing  $\log p(x_n)$  directly. Because:

$$p(x) = \int p(z)p(x | z)dz$$

In  $p(x)$ , the integral is intractable, especially when  $z$  has high dimension.

#### Question 1.5

$$\log p(x_n) - KL(q(Z | x_n) || p(Z | x_n)) = ELBO$$

When the lower bound is pushed up, two things can happen:

- 1) Log likelihood  $\log p(x_n)$  is maximized.
- 2)  $KL(q(Z | x_n) || p(Z | x_n))$  is minimized. This means that the KL divergence between approximated posterior and true posterior become smaller. The learned posterior approaches to the true posterior. Ideally, when the learned latent variable is perfectly equal to the true latent variable,  $KL(q(Z | x_n) || p(Z | x_n)) = 0$ ,  $\log p(x_n) = ELBO$

#### Question 1.6

Reconstruction loss:  $L_n^{rec} = -E_{q_\phi(z|x_n)}[\log p(x_n | z)]$  it computes the negative log likelihood of the reconstructed data, with the given latent variable. In the equation,  $q_\phi(z | x_n)$  is the approximated posterior. We firstly sample  $z$  from this approximated posterior. Then we reconstruct  $x_n$ . The log likelihood  $\log p_\theta(x_n | z)$  assigns high probability to the reconstruction. When minimizing the loss function, the smaller the loss  $L_n^{recon}$  is, the larger the log likelihood  $E_{q_\phi(z|x_n)}[\log p_\theta(x_n | z)]$  is. We want the likelihood  $p_\theta(x_n | z)$  to be as large as possible, which means the reconstructed data  $x_n$  has a high possibility to be generated from  $z$ .

Regularization loss:  $L_n^{rec} = D_{KL}(q_\phi(z | x_n) || p_\theta(z))$  it computes the KL divergence (dissimilarity) between the approximated posterior  $q_\phi(z | x_n)$  and the prior  $p_\theta(z)$ . We want  $L_n^{rec}$  to be as small as possible. It means that we force the approximated posterior (or the encoder)  $q_\phi(z | x_n)$  to be similar as the prior. The regularization loss is used to avoid overfitting.

#### Question 1.7

$$\begin{aligned}
L_n^{\text{recon}} &= -E_{q_\phi(z|x_n)}[\log(p_\theta(x_n | z))] \\
&= -\log(p_\theta(x_n | z)) \\
&= -\log \prod_{m=1}^M (\text{Bern}(x_n^{(m)} | f_\theta(z_n)_m)) \\
&= -\sum_{m=1}^M \log(\text{Bern}(x_n^{(m)} | f_\theta(z_n)_m)) \\
&= -\sum_{m=1}^M \log(f_\theta(z_n)_m^{x_n^{(m)}} \bullet (1 - f_\theta(z_n)_m)^{1-x_n^{(m)}}) \\
&= -\sum_{m=1}^M \{\log(f_\theta(z_n)_m^{x_n^{(m)}}) + \log(1 - f_\theta(z_n)_m)^{1-x_n^{(m)}}\} \\
&= -\sum_{m=1}^M \{x_n^{(m)} \log(f_\theta(z_n)_m) + (1 - x_n^{(m)}) \log(1 - f_\theta(z_n)_m)\}
\end{aligned}$$

$$\begin{aligned}
L_n^{\text{reg}} &= D_{KL}(q_\phi(z | x_n) \| p_\theta(z)) \\
&= D_{KL}(N(z_n | \mu_\phi(x_n), \text{diag}(\Sigma_\phi(x_n))) \| N(z_n | 0, I_D)) \\
&= \sum_d^D \frac{1}{2} \{\text{var}_q + \text{mean}_q^2 - 1 - \log(\text{var}_q)\} \\
&= \frac{1}{2} \sum_d^D \{\text{diag}(\Sigma_\phi(x_n)) + \mu_\phi(x_n)^T \mu_\phi(x_n) - 1 - \log(\text{diag}(\Sigma_\phi(x_n)))\} \\
&= \frac{1}{2} \left\{ \sum_d^D \text{diag}(\Sigma_\phi(x_n)) + \sum_d^D \mu_\phi(x_n)^T \mu_\phi(x_n) - \sum_d^D 1 - \sum_d^D \log(\text{diag}(\Sigma_\phi(x_n))) \right\} \\
&= \frac{1}{2} \left\{ \text{Tr}(\text{diag}(\Sigma_\phi(x_n))) + \sum_d^D \mu_\phi(x_n)^T \mu_\phi(x_n) - D - \sum_d^D \log(\text{diag}(\Sigma_\phi(x_n))) \right\}
\end{aligned}$$

$$\begin{aligned}
L &= \frac{1}{N} \sum_{n=1}^N (L_n^{\text{recon}} + L_n^{\text{reg}}) \\
&= \frac{1}{N} \sum_{n=1}^N \left\{ -\sum_{m=1}^M (x_n^{(m)} \log(f_\theta(z_n)_m) + (1 - x_n^{(m)}) \log(1 - f_\theta(z_n)_m)) + \right. \\
&\quad \left. \frac{1}{2} \sum_d^D (\text{diag}(\Sigma_\phi(x_n)) + \mu_\phi(x_n)^T \mu_\phi(x_n) - 1 - \log(\text{diag}(\Sigma_\phi(x_n)))) \right\}
\end{aligned}$$

*D: dimension of latent space*

*N: total amount of images*

*M: total amount of pixels in each image*

*$x_n$ : the  $n^{\text{th}}$  image*

#### Question 1.9 Reparameterization trick

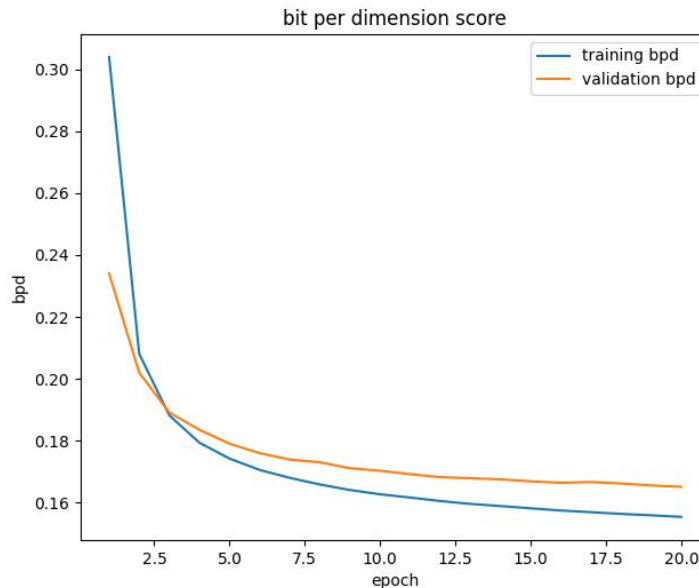
The problem is that we want to compute the gradient of loss with respect to  $\theta$  and  $\phi$ . However, it is impossible to compute the gradient. Because sampling process in the forward pass is stochastic. So it is not differentiable. Therefore, we use the reparameterization trick.

Reparameterization trick is: we add noise to the standard deviation of the latent variable. The noise is drawn from normal distribution  $N(0,1)$ . After adding noise, latent variable  $z$  is written as:

$$\begin{aligned}
Z &= \mu_z + \epsilon \sigma_z \\
\epsilon &\in N(0,1), Z \in N(\mu_z, \sigma_z)
\end{aligned}$$

Reparameterization trick allows us to compute the gradient in back propagation. Because it disconnects the random sampling by using the external noise epsilon. In backpropagation, the gradients flow through deterministic functions: mean and standard deviation of  $z$ . Randomness has been moved to the noise.

### Question 1.9



### Architecture

Encoder: its first layer is Linear (in=784, out=512). Its second layer is Relu. Based on this output, it splits into two layers: one linear layer for mean and another linear layer for log variance. The mean and log var both have the same shape: batch size \* dimensionality of latent space.

Decoder: It has a linear layer (in=z\_dim, out=512), a relu, and a linear layer (in=512, out=784). Output of decoder is reconstructed images which has shape: batch size \* channel \* H \* W

### Question 2.1

A)  $E_{p_{data}(x)}[\log D(X)]$  is the expectation of log probability that discriminator classifying  $x$ .  $x$  is drawn from the distribution of real data. If discriminator classifies real data as real, then  $D(x)=1$ .

$E_{p_z(x)}[\log(1-D(G(X)))]$  is the expectation of log probability that the discriminator classifying the fake  $x$ . fake  $x$  is  $G(x)$ , which is generated by the

generator function  $G$ . If discriminator classifies fake data as real, then  $D(G(x))=1$ . If discriminator classifies fake data as fake, then  $D(G(x))=0$ .


B) After training has converged: the learned generator distribution is identical to  $p_{\text{data}}$ .  $p_{\text{data}} = p_g$ . So the optimal discriminator is  $1/2$ . Discriminator and generator reaches equilibrium, when training is converged.

$$D^{\text{optimal}}(x) = \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_g(x)} = \frac{1}{2}$$

$$V(D^{\text{optimal}}, G) = \log\left(\frac{1}{2}\right) + \log\left(1 - \frac{1}{2}\right) = -\log 4$$

## Question 2.2

Problem: When sample is likely fake, we want to learn from it to improve the generator. However,  $\log(1-D(G(z)))$  is too flat in the region where  $D(G(z))$  is close to 0. It means that  $\log(1-D(G(z)))$  has low gradient signal when the discriminator classifies data as fake.



Solution: Instead of minimizing likelihood of discriminator being correct, we maximize likelihood of discriminator being wrong. We assign higher gradient signal for bad samples. We convert  $\log(1-D(G(z)))$  to  $\log(D(G(z)))$ :

$$\min E[\log(1 - D(G(z)))] \Leftrightarrow \max E[\log(D(G(z)))]$$

Now the gradient has high signal in the region where  $D(G(z))$  is close to zero.

## Question 2.3

### Architecture:

generator	discriminator
It contains 4 hidden modules. Each module contains: a linear layer, a batchnorm layer, a leakyrelu layer and a dropout layer. In the end, it has a linear layer and a tanh layer.	It contains 2 modules. Each module contains: a linear layer, a leakyrelu layer and a dropout layer. In the end, it has a linear layer.

### Question 2.5

One problem is: Batch-normalization causes strong intra-batch correlation. A lot of samples generated within the batch tend to be similar. Because the activation between different minibatches depends on each other. So the generations (product of activation) depends on each other. Solution: virtual batch normalization. We select a reference batch which is fixed during training. We append the reference batch to regular mini-batch. We combine it with the statistics of the current batch.

### Question 3.1



$p_x(x) = p(z) \left  \frac{dz}{dx} \right $ $= \frac{1}{b-a} \left  \frac{d}{dx} f(x) \right $ $= \frac{1}{b-a} \left  \frac{d}{dx} x^3 \right $ $= \frac{1}{b-a} 3x^2$ $z = f(x) = x^3 \Rightarrow x = f(z)^{-1} = z^{\frac{1}{3}}$ $z \in (a, b) \Rightarrow x \in (a^{\frac{1}{3}}, b^{\frac{1}{3}})$	$\text{verify: } \int_{-\infty}^{\infty} p_x(x) dx = 1$ $\int_{-\infty}^{\infty} p_x(x) dx$ $= \int_{-\infty}^{\infty} \frac{1}{b-a} 3x^2 dx$ $= \frac{3}{b-a} \int_{-\infty}^{\infty} x^2 dx$ $= \frac{3}{b-a} \int_{a^{\frac{1}{3}}}^{b^{\frac{1}{3}}} x^2 dx$ $= \frac{3}{b-a} \left( \frac{1}{3} x^3 \Big _{x=b^{\frac{1}{3}}} - \frac{1}{3} x^3 \Big _{x=a^{\frac{1}{3}}} \right)$ $= \frac{3}{b-a} \left( \frac{1}{3} (b^{\frac{1}{3}})^3 - \frac{1}{3} (a^{\frac{1}{3}})^3 \right)$ $= \frac{1}{b-a} (b-a)$ $= 1$
---	--

### 3.2

a)  $h_l$  must be 1) differentiable 2) invertible 3) smooth 4) easy to compute the inverse of  $h_l$  5) easy to compute the jacobian determinant . Because


when the number of dimension in  $h_l$  and  $h_{l-1}$  is high,  $\left| \frac{dh_l}{dh_{l-1}} \right|$  will have high

dimension as well. Otherwise, it is hard to compute  $\log p_x(x)$  .


b) The first optimization issue: very likely,  $h_l$  can not satisfy all the above requirements. The second optimization issue: in equation 18, inverse of a function and jacobian determinant are very hard to compute.

### 3.3

A) When variables are discrete, like images, instead of continuous, the learned distribution looks like the following image. The black points represent the discrete points, and the green volume the density modeled by a normalizing flow in continuous space. For  $x=0,1,2,3$ , they have infinite high likelihood. Because their interval is 0 in the continuous space. Because  $\int p(x)dx=1$ . The learned density does give any useful information about the distribution among the discrete points.



To fix it, we use dequantization. Dequantized  $v = x + u$ , where  $u \in [0,1)^D$ . It makes the probability of  $x$  to become smooth. As shown in the following image, the shape of  $p(x=0)$  is smoothly connected with  $p(x=1)$ .



c) the steps to train a general flow based model, evaluate it after training, sample from it.

Step 1 assume a distribution function for  $p(z)$

Step 2 go through forward pass, from  $x$  to  $z$

Step 3 compute loss between the  $p(z)$  that we want and the  $p(z)$  that the model computes

Step 4 go backward and update the parameters of  $f(x)$ .

Step 5 repeat step 3 and 4 until the  $p(z)$  reaches our goal

Step 6 now we are able to compute the likelihood  $p(x)$ . because we have learned the function  $f(x)$ . we go back from  $z$  to  $x$ , which is the inverse of function  $f$ , then we compute  $p(x)$

## Conclusion


GAN uses minimax game to maximize the likelihood. Generator aims at fooling the discriminator by generating data which are fake but looks real. Discriminator aims at correctly distinguishing real and fake data. Latent variable is the source of noise.

VAE uses lower bound to maximize the likelihood. VAE uses latent variable  $z$  to capture the feature of input variables. Then it reconstructs new data based on  $z$ . It encode  $x$  to  $z$ , then it decode  $z$  to  $x$ . It is like doing zipping. The dimension is reduced through encoder.


Flow model: Each input data has a corresponding latent variable. The dimension of  $z$  is the same as the dimension of  $x$ .

VAE computes an approximated likelihood. Flow Model computes the exact likelihood directly. GAN computes implicit density while VAE computes explicit density.

**GAN:** minimax the classification error loss.



**VAE:** maximize ELBO.



**Flow-based generative models:** minimize the negative log-likelihood

